

edX

Course

Discussion

Wiki

Progress

Course > Bandits > Lab > Exercise 2A Epsilon ...

<

>

# Exercise 2A Epsilon Greedy

Bookmark this page

## Exercise 2.2A: Epsilon Greedy

In this exercise, you will examine the epsilon greedy policy.

Make sure that you have:

- Completed the setup requirements as described in the Set Up Lab Environments section
- Completed the previous exercises in this lab

Now, run jupyter notebook and open the “Ex2.2A Epsilon Greedy.ipynb” notebook under **Module 2** folder.

- Examine the notebook. We have given you an implementation of the epsilon greedy algorithm.
- Once you have studied the notebook, prepare a simulation. Don't change any other parameter, that is:
  - evaluation\_seed = 5016
  - num\_actions = 10
  - trials = 10000
  - distribution = “bernoulli”
- Set the epsilon to zero.
- Run the simulation, observe the results, and answer the following questions.

## Lab Question

1/1 point (graded)

With epsilon set to zero, what do you observe?

☐ The epsilon greedy behaves randomly

☐ The epsilon greedy behaves like the optimistic greedy algorithm

☒ The epsilon greedy behaves like the greedy algorithm

Submit

You have used 1 of 2 attempts

Save

Show Answer

Correct (1/1 point)

Set the epsilon to one. Run the simulation again and observe the results.

## Lab Question

1/1 point (graded)

With epsilon set to one, what do you observe?

☒ The epsilon greedy behaves randomly

☐ The epsilon greedy behaves like the optimistic greedy algorithm

☐ The epsilon greedy behaves like the greedy algorithm

Submit

You have used 1 of 2 attempts

Save

Show Answer

Correct (1/1 point)

Now, try several different number of epsilons (0.05, 0.1, 0.15). Make sure the other parameters stay the same, that is:

- evaluation\_seed = 5016
- num\_actions = 10
- trials = 10000
- distribution = “bernoulli”

Run the simulations and observe the results.

## Lab Question

1/1 point (graded)

Which epsilon gives the most average reward?

☒ 0.05

☐ 0.1

☐ 0.15

Submit

You have used 1 of 2 attempts

Save

Show Answer

Correct (1/1 point)

Now let's prepare another simulation by setting a different seed, so your parameters should look like this:

- evaluation\_seed = 1239
- num\_actions = 10
- trials = 10000
- distribution = “bernoulli”

Run the simulations with different number of epsilons (0.05, 0.1, 0.15) and observe the results.

## Lab Question

1/1 point (graded)

Which epsilon gives the most average reward?

☐ 0.05

☒ 0.1

☐ 0.15

Submit

You have used 1 of 2 attempts

Save

Show Answer

Correct (1/1 point)

< Previous

Next >